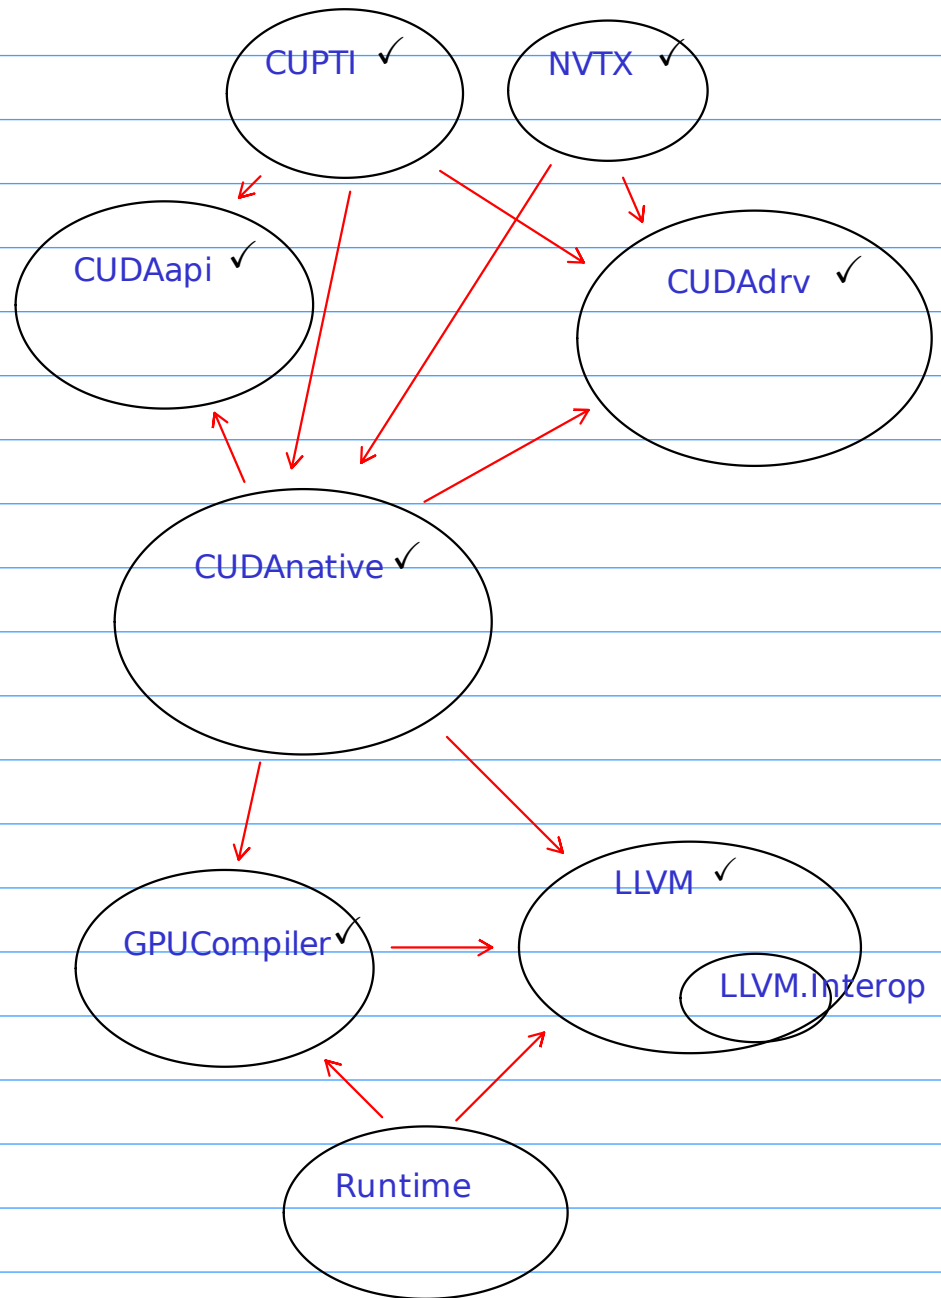


CUDAnative.jl

CUDAnative.jl

CUPTI: CUDA Profiling Tools Interface
NVTX: NVIDIA Tools Extension SDK

GPUのperf counter群を扱うhigh-level API?
annotationして超高級printf debugみたいな?



module 相関図

Why LLVM.jl

The high level Julia IR is accessible with the `code_lowered` and `code_typed` reflection functions, and can be modified with generated functions.

Low-level LLVM IR can be inspected by invoking `code_llvm` and injected via the `llvmcall` metaprogramming interface. Machine code is accessible through `code_native` and can be inserted indirectly as inline assembly in LLVM IR.

LLVM IR への interface は string based

例えばpointerを指定してloadする関数をstring-basedで書いたらその関数はstring操作で溢れて読めたもんじゃないし各pointer typeを個別に扱ってoptimizeするようにはめに陥るだろう

```
julia> using LLVM
julia> mod = LLVM.Module("llvmcall")
; ModuleID = 'llvmcall'
source_filename = "llvmcall"
julia> eltyp = LLVM.convert(LLVMType,Int64)
i64
julia> pram_typs = [LLVM.PointerType(eltyp)]
1-element Array{LLVM.PointerType,1}:
i64*
julia> ft = LLVM.FunctionType(eltyp, pram_typs)
i64 (i64*)
```

```
julia> LLVM.Builder() do builder
    bb = LLVM.BasicBlock(f, "entry")
    LLVM.position!(builder, bb)
    ptr = LLVM.parameters(f)[1]
    val = LLVM.load!(builder, ptr)
    LLVM.ret!(builder, val)
end
ret i64 %1
```

```
julia> f
```

```
define i64 @load(i64*) {
entry:
    %1 = load i64, i64* %0
    ret i64 %1
}
```

LLVM.jl

lib/

@apicall(:LLVMなんたらかんたら , ...) C function?

Clang.jlで生成?

wrap/wrap.jl

llvm-configを走らせてheaderを生成したりする?

src

almost except LLVM.jl are LLVM API wrappers

- └ LLVM.jl
- └ analysis.jl
- └ base.jl
- └ bitcode.jl
- └ buffer.jl
- └ core
- └ core.jl
- └ datalayout.jl
- └ debuginfo.jl
- └ deprecated.jl
- └ execution.jl
- └ init.jl
- └ interop
- └ interop.jl
- └ ir.jl
- └ irbuilder.jl
- └ linker.jl
- └ moduleprovider.jl
- └ pass.jl
- └ passmanager.jl
- └ passregistry.jl
- └ support.jl
- └ target.jl
- └ targetmachine.jl
- └ transform.jl
- └ types.jl
- └ util

base: Define @apicall

```
libllvm_paths = filter(Libdl.dllist()) do lib
    occursin("LLVM", basename(lib))
end
```

filter ideom

new{T,S} return something{T,S} type object

GPUCompiler.jl

@kwdef typedef
automatically defines a keyword-based constructor

ptx.jl hide_trap!
pass removes calls to `trap` and replaces them with inline assembly

CUDAnative.jl

CUDAnative.jl のすること:
CUDAのconfigurationをチェック
device/*.jl を include
CUDACompilerTarget, CUDACompilerJob を定義
*.jl を include
initialization
__init__(), __configure__(), __runtime_init__()

device/tools.jl Create MDNodes
device/pointer.jl Pointers with address space information, AS module, DevicePtr...
Generate LLVM IR for memory operation
device/array.jl Contiguous on-device arrays
device/cuda.jl Include device/cuda/*.jl
device/llvm.jl wrappers for LLVM-specific functionality: trap(), assume()

?bitcast

init.jl CUDA Initialization and Context Management
context, context!, device!, device_reset!
CUDAnative.attaskswitch, CUDAnative.atcontextswitch
compatibility.jl Julia, CUDA and LLVM version compatibility
bindeps.jl CUDA discovery and dependency

cuoti/CUPTI.jl Wrappers for CUPTI and NVTX
nvtx/NVTX.jl

execution.jl @cuda
exceptions.jl device-side exceptions
reflection.jl @device_code_*
array.jl CUDA host array type for testing purposes

deprecated.jl

device/cuda/*.jl

memory_shared.jl

@cuStaticSharedMem(T::Type, dims) -> CuDeviceArray{T,AS.Shared}

@cuDynamicSharedMem(T::Type, dims, offset::Integer=0) -> CuDeviceArray{T,AS.S

These 2 call _shmem @generated function which creates llvm function call so to returns a pointer to shared memory.

indexing.jl

threadIdx, blockDim, blockIdx, gridDim, warpsize

llvm.nvvm.read.ptx.sreg.tid.x などと呼ぶ generated function で

llvm.nvvm.read.ptx.sreg.warpsize を ccall

ATM, values are based on V100

synchronization.jl

sync_threads, sync_warp,

sync_threads_count, sync_threads_and, sync_threads_or

threadfence, threadfence_block, threadfence_system

llvm.nvvm.* function を ccall

warp_vote.jl

vote_all, vote_any, vote_ballot

@asmcallでPTXのvote insnを呼び出し

warp_shuffle.jl

llvm.nvvm.shfl.sync.\$mode.\$typ intrinsicをccallする関数群を生成

output.j

@cuprintf

Print a formatted string in device context on the host standard output.

@cuprint, @cuprintln

Print a textual representation of args to standard output from the GPU.

\$(map(esc, args)...)

asserstion.jl

@cuassert

@assert cond [text]

[continued]

memory_dynamic.jl
malloc

atomics.jl
##LLVM
@generated function llvm_atomic_op(::Val{binop}, ptr::DevicePtr{T,A}, val::T)
 where {binop, T, A}
 arithmetic operations on integers using LLVM instructions
@generated function llvm_atomic_cas(ptr::DevicePtr{T,A}, cmp::T, val::T)
 where {T, A}
##NVVM
@llvm.nvvm.atomic.load.add.*
@llvm.nvvm.atomic.load.inc.*
@llvm.nvvm.atomic.load.dec.*
##Julia
atomic_*!()
Higher level interface
@atomic a[l] = op(a[l], val)
@atomic a[l] ...= val

misc.jl
clock, nanosleep
llvm.nvvm.read.ptx.sreg.clock64 and nanosleep.u32 insn

wmma.jl
Warp Matrix Multiply Accumulate ???

" The libdevice library is a collection of NVVM bitcode functions that implement common functions for NVIDIA GPU devices, including math primitives and bit-manipulation functions. These functions are optimized for particular GPU architectures, and are intended to be linked with an NVVM IR module during compilation to PTX."

math.jl

" The libcudadevrt library is a collection of PTX bitcode functions that implement parts of the CUDA API for execution on the device, such as device synchronization primitives, dynamic kernel APIs, etc."

libcudadevrt_common.jl, libvudadevrt.jl : cuda_device_runtime_api

cooperative_groups.jl
this_grid, sync_grid

dynamic_parallelism.jl
launch, synchronize

CUDAdrv.jl

CUDAdrv.jl のすること

pointer.jl が essential functionality

CuPtr{T}, PtrOrCuPtr{T}

各種 wrapper を include する

Do init that can't be done at module load time

deferred initialization

initialization

__configure__() which is essentially ccall of cuInit

low-level wrappers

libcuda_common.jl

error.jl CuError, CUDAdrv.initializer(f::Function)

libcuda.jl wrapper gened from cuda.h

libcuda_aliases.jl Remove _v2 postfix

high-level wrappers

version.jl version() and release() of CUDA

devices.jl CuDevice, name, totalmem, attribute,...

context.jl CuContext,...

context/primary.jl CuPrimaryContext

stream.jl CuStream,...

memory.jl Mem, Mem.DeviceBuffer, Mem.Device,
Mem.alloc, Mem.[un]register, Mem.free, copy...

module.jl include module/jit.jl encode/decode JIT options

CuModule, CuModuleFile

include module/function.jl CuFunction

include module/global.jl CuGlobal, get, set

include module/linker.jl CuLink,...

events.jl CuEvent, record, synchronize, elapsed

execution.jl CuDim, cudacall

cudacall calls launch() with convert_argument

profile.jl profiling with Nsight/NVTX or with external profiler

occupancy.jl active_blocks, occupancy, launch_configuration

deprecated.jl

cuXXXはCUDA_driver, cudaXXXはCUDA_runtime

CUDAapi.jl

CUDAのinstall状況を知る+お役立ちmodule

CUDAapi.jl includes

util.jl	@enum_without_prefix
discovery.jl	find_cuda_library, find_cuda_binary, find_toolkit, parse_toolkit_version, find_libdevice, find_libcudadevrt
availability.jl	has_cuda, has_cuda_gpu, usable_cuda_gpus call find_toolkit and ccall cuDeviceGetCount
complex.jl	Define cuda complex types
library_types.jl	Other cuda data types
call.jl	@runtime_ccall, decode_ccall_function, @checked ccall extention wrappers

GPUCompiler.jl Take2

GPUCompiler.jl includes *.jl below

utils.jl tbaa_make_child child用LLVM metadata nodeを返す? **TBAA???**

compiler interface and implementations
interface.jl target-specific interface for packages to implement
FunctionSpec function, kernelのaggregate type

CompilerTarget, CompilerJob abstractな操作

error.jl error handling
KernelError, InternalCompilerError

native.jl native target(Sys.MACHINE) for CPU execution
NativeCompilerTarget, NativeCompilerJob

ptx.jl PTX target(nvptx64-nvidia-cuda) for GPU execution
PTXCompilerTarget, PTXCompilerJob

ここは大変

gcn.jl GCN target(amdgcen-amd-amdhsa) for GPU execution
GCNCompilerTarget, GCNCompilerJob

runtime.jl GPU のruntimeで必要なJulia runtime functions [module Runtile]
それらの関数はRuntimeMethodInstanceとして保持される. もしその関数が
targetで規定されるものならUndefVarErrorを避けるためllvmcall stubを生成する
malloc, box, unboxなどのsupportもここで

compiler implementation
irgen.jl LLVM IR generation
compile_method_instance LLVM IRを生成しLLVM moduleを作って返す
irgen さらにtarget-specific processingやentryの処理などを行う
promote_kernel! promote a function to a kernel

optim.jl LLVM IR optimization, kernel lowering
validation.jl Validate method, IR, ...
rtlib.jl Compiler support for GPU runtime library
mcgen.jl final preparations for the module to be compiled to PTX
debug.jl backtrace
driver.jl Compiler driver and main interface: compile() which calls
codegen() which makes passes run これが大概

other reusable functionality
cache.jl Compilation cache
execution.jl Helper to implement code exec. split_kwargs, assign_args!
reflection.jl code_*, device_code_*